

LAB CYCLE 4

Experiment No:4

Date:10/02/22

Aim:

Convert all the values in quality attribute to 0 (bad) if the value is less than '5', to 1 (good) if the value is '5' or '6' and to 2 (great) otherwise. Normalize all the other attributes by Z-score normalization, and segregate them into 4 equal spaced bins each giving the values between [0 to 3], and replace the values for that attribute with the number corresponding to the interval they belong.

For example, suppose after normalization an attribute has values between [-0.5,1.5], i.e., minimum value of the attribute is -0.5 and maximum value is 1.5, then form 4 bins:

bin 0: [-0.5,0.0],

bin 1: [0.0,0.5],

bin 2: [0.5,1.0],

bin 3: [1.0,1.5].

For example, if a data instance has a value of 0.73 for that attribute, replace 0.73 with 2. Use this dataset for constructing a Decision Tree.

Problem Statement

1. Implement Decision tree algorithm using information gain to choose which attribute to split at each point. Stop splitting a node if it has less than 10 data points. Do NOT use scikit-learn for this part.
2. Test out the implementation of Decision Tree Classifier from scikit-learn package, Using information gain. Here also stop splitting a node if it has less than 10 data points.
3. Cross validate the classifiers with 3-folds and print the mean macro accuracy, macro precision and macro recall for both the classifiers. You may or may not use the scikit-learn implementations for computing these metrics and cross validation.

Source Code:

```
In [1]: import numpy as np
import pandas as pd
from pprint import pprint
from sklearn.tree import DecisionTreeClassifier
```

```
In [2]: dt=pd.read_csv("winequality-red.csv")
```

```
In [3]: dt
```

```
Out[3]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	9.8	5
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	9.8	5
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	9.8	6
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	9.4	5
...
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	10.5	5
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	11.2	6
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	11.0	6
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	10.2	5
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	11.0	6

1599 rows × 12 columns

```
In [4]: n=dt.shape
n
```

```
Out[4]: (1599, 12)
```

```
In [51]: dt.head(8)
```

```
Out[51]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	5
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	5
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	6
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	5
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4	5
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4	5
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0	7

```
In [52]: dt.loc[dt["quality"]<5,"quality"]=0
dt.loc[dt["quality"]==5,"quality"]=1
dt.loc[dt["quality"]==6,"quality"]=1
dt.loc[dt["quality"]<6,"quality"]=2
```

```
In [53]: dt.head(8)
```

```
Out[53]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	2
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8	2
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8	2
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8	2
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4	2
5	7.4	0.66	0.00	1.8	0.075	13.0	40.0	0.9978	3.51	0.56	9.4	2
6	7.9	0.60	0.06	1.6	0.069	15.0	59.0	0.9964	3.30	0.46	9.4	2
7	7.3	0.65	0.00	1.2	0.065	15.0	21.0	0.9946	3.39	0.47	10.0	7

```
In [54]:
def normalize(x):
    xnew=((x-np.mean(x))/np.std(x))

    #print(xnew)
    return xnew
```

```
In [55]: dt.iloc[:,0:11]=dt.iloc[:,0:11].apply(normalize)
```

```
In [56]: dt.head(8)
```

```
Out[56]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	-0.528360	0.961877	-1.391472	-0.453218	-0.243707	-0.466193	-0.379133	0.558274	1.288643	-0.579207	-0.960246	2
1	-0.298547	1.967442	-1.391472	0.043416	0.223875	0.872638	0.624363	0.028261	-0.719933	0.128950	-0.584777	2
2	-0.298547	1.297065	-1.186070	-0.169427	0.096353	-0.083669	0.229047	0.134264	-0.331177	-0.048089	-0.584777	2
3	1.654856	-1.384443	1.484154	-0.453218	-0.264960	0.107592	0.411500	0.664277	-0.979104	-0.461180	-0.584777	2
4	-0.528360	0.961877	-1.391472	-0.453218	-0.243707	-0.466193	-0.379133	0.558274	1.288643	-0.579207	-0.960246	2
5	-0.528360	0.738418	-1.391472	-0.524166	-0.264960	-0.274931	-0.196679	0.558274	1.288643	-0.579207	-0.960246	2
6	-0.241094	0.403229	-1.083370	-0.666062	-0.392483	-0.083669	0.381091	-0.183745	-0.072005	-1.169337	-0.960246	2
7	-0.585813	0.682553	-1.391472	-0.949853	-0.477498	-0.083669	-0.774449	-1.137769	0.511130	-1.110324	-0.397043	7

```
In [57]:
bin_labels_4=[0,1,2,3]
#dt['fixed acidity'] = pd.qcut(dt['fixed acidity'],q=4,label=bin_labels_4)
for i in range(0,11):
    dt.iloc[:,i]=pd.qcut(dt.iloc[:,i],q=4,labels=bin_labels_4)
dt.head(8)
```

```
Out[57]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	1	3	0	0	1	1	1	2	3	1	0	2
1	1	3	0	2	3	3	3	2	0	2	1	2
2	1	3	0	2	3	2	2	2	1	2	1	2
3	3	0	3	0	1	2	2	3	0	1	1	2
4	1	3	0	0	1	1	1	2	3	1	0	2
5	1	3	0	0	1	1	2	2	3	1	0	2
6	1	2	0	0	0	2	2	1	1	0	0	2
7	1	3	0	0	0	2	0	0	2	0	1	7

```
In [58]:
traincount=int(dt.shape[0]*0.8)
traincount
```

Out[58]: 1279

```
In [59]: def train_test_split(dt):
         training_data=dt.iloc[:traincount].reset_index(drop=True)
         testing_data=dt.iloc[traincount:].reset_index(drop=True)
         return training_data,testing_data
         training_data=train_test_split(dt)[0]
         testing_data=train_test_split(dt)[1]
```

```
In [60]: training_data.shape
```

Out[60]: (1279, 12)

```
In [61]: testing_data.shape
```

Out[61]: (320, 12)

```
In [62]: #Compute entropy
```

```
In [63]: def entropy(class_label):
         values,counts=np.unique(class_label,return_counts=True)

         for i in range(len(values)):
             entropy=np.sum([(count[i]/np.sum(counts))*np.log2(counts[i]/np.sum(counts))])
         return entropy
```

```
In [64]: #Info Gain
```

```
In [65]: def InfoGain(data,split_attribute_name,class_label="equality"):
         total_entropy=entropy(data[class_label])
         vals,counts=np.unique(data[split_attribute_name],return_counts=True)
         #Calculate the weighted entropy
         for i in range(len(vals)):
             Weighted_Entropy=np.sum([(counts[i]/np.sum(counts))*entropy(data.where(data[split_attribute_name]==va

         #formula for information gain
         Information_Gain=total_entropy-Weighted_Entropy
         return Information_Gain
```

```
In [66]: def ID3(data,originaldata,features,class_label="quality",parent_node_class=None):
         #if all class_label values are same, return that value
         if len(np.unique(data[class_label]))<=1:
             return np.unique(data[class_label])[0]

         #if the dataset is empty or below some threshold value,terminate recursion
         elif len(data)==0:
             #find the counts of distinct values of class_label, then find the maximum count of them--> majority c
             return np.unique(originaldata[class_label])[np.argmax(np.unique(originaldata[class_label],return_coun

         #if the feature space is empty,terminate recursion
         elif len(features)==0:
             return parent_node_class

         #If none of the above condition holds true form the subtrees

         else:
             #Find the counts of distinct values of class_label, then find the maximum count of them-->majority cl
             parent_node_class=np.unique(data[class_label])[np.argmax(np.unique(data[class_label],return_counts=Tr

         #Select the feature which best splits the dataset, feature having maximum information gain

         for feature in features:
             item_values=[InfoGain(data,feature,class_label)] #Return the infogain values
             best_feature_index=np.argmax(item_values)
             best_feature=features[best_feature_index]

         #Create the tree structure as a nested dictionary
```

```

tree={best_feature:{}}

#Remove the feature with the best info gain
features=[i for i in features if i!=best_feature]

#Form subtrees down the root node by calling ID3 recursively

for value in np.unique(data[best_feature]):
    value=value
    sub_data=data.where(data[best_feature]==value).dropna()
    #call the ID3 algorith
    subtree=ID3(sub_data,dt,features,class_label,parent_node_class)
    #Add the subtree
    tree[best_feature][value]=subtree
return(tree)

```

In [67]:

```

tree = ID3(training_data,training_data,training_data.columns[:-1])
pprint(tree)

```

```

{'fixed acidity': {0: {'volatile acidity': {0: {'citric acid': {0: 2.0,
1: 2.0,
2: {'residual sugar': {0: {'chlorides': {0:
{'free sulfur dioxide': {0: 7.0,
2: 2.0,
3: 2.0}},
2.0,
2.0}},
1:
2:
2.0}},
1: 2.0,
2: {'chlorides': {0:
3:
3: 2.0}},
3: {'residual sugar': {0: 2.0,
1: 7.0,
2: 2.0}}}},
1: {'citric acid': {0: {'residual sugar': {0: {'chlorides': {0:
{'free sulfur dioxide': {2: {'total sulfur dioxide': {1: 2.0,
2: 2.0,
3: 7.0}},
3: {'total sulfur dioxide': {1: 2.0,
3: 8.0}}}},
1:
2:
2.0}},
1: {'chlorides': {0:
{'free sulfur dioxide': {0: 2.0,
1: 2.0,
2: 7.0}},
2.0,
2.0}},
2: {'chlorides': {0:
{'free sulfur dioxide': {1: 7.0,
2: 2.0}},
2.0}}}},
1:
1: {'residual sugar': {0: 2.0,
1: {'chlorides': {0:
{'free sulfur dioxide': {0: 2.0,

```

```

2.0,
2.0,
2.0,
{'free sulfur dioxide': {1: 2.0,
3: 7.0}}}},
2.0,
{'free sulfur dioxide': {0: {'total sulfur dioxide': {0: 8.0,
1: 2.0,
2: 7.0}},
1: 2.0}},
{'free sulfur dioxide': {0: {'total sulfur dioxide': {0: {'density': {2: 7.0,
3: {'pH': {0: 7.0,
1: 2.0}}}}}},
1: 2.0,
3: 2.0}}}}}}}},
3: {'citric acid': {1: 2.0,
2: {'residual sugar': {0: 2.0,
1: 2.0,
2: 2.0,
3: {'chlorides': {1:
2.0,
7.0,
2.0}}}},
3: {'residual sugar': {0: 2.0,
2: 2.0,
3: {'chlorides': {1:
2.0,
2.0,
{'free sulfur dioxide': {0: 7.0,
1: 2.0}}}}}}}}}}}}

```

```

In [68]: def predict(query,tree,default = 1):
#1.
for key in list(query.keys()):
    if key in list(tree.keys()):
        #2.
        try:
            result = tree[key][query[key]]
        except:
            return default

        #3.
        result = tree[key][query[key]]
        #4.
        if isinstance(result,dict):
            return predict(query,result)

    else:
        return result

```

```

In [69]: def test(data,tree):
#Create new query instances by simply removing the target feature column from the original dataset and
#convert it to a dictionary
queries = data.iloc[:, :-1].to_dict(orient = "records")

#Create a empty DataFrame in whose columns the prediction of the tree are stored

```

```

predicted = pd.DataFrame(columns=["predicted"])

#Calculate the prediction accuracy
for i in range(len(data)):
    predicted.loc[i,"predicted"] = predict(queries[i],tree,1.0)
print('The prediction accuracy is: ',(np.sum(predicted["predicted"] == data["quality"])/len(data))*100,'%')

```

In [70]: `test(testing_data,tree)`

The prediction accuracy is: 78.75 %

In []:

In [72]: `#Decision tree implemenentation using Libraries`

```

from sklearn.model_selection import train_test_split#or decision rree object
from sklearn.tree import DecisionTreeClassifier#or checking testing resutts
from sklearn.metrics import classification_report

```

In [73]: `#to divide data into attributes and labels,execute the following code:`

```

X = dt.drop('quality', axis=1)
y = dt['quality']

```

In [74]: `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)`

In [75]: `#TTraining and Making Predictions`

```

classifier = DecisionTreeClassifier()
classifier.fit(X_train, y_train)

```

Out[75]: `DecisionTreeClassifier()`

In [76]: `y_pred = classifier.predict(X_test)`

In [77]: `print(classification_report(y_test, y_pred))`

	precision	recall	f1-score	support
2	0.93	0.92	0.92	282
7	0.39	0.41	0.40	34
8	0.00	0.00	0.00	4
accuracy			0.85	320
macro avg	0.44	0.44	0.44	320
weighted avg	0.86	0.85	0.85	320

In [37]: `target = list(dt['quality'].unique())`
`feature_names = list(X.columns)`
#We can also get a textual representation of the tree by using the export tree function from the SkLearn Libr
`from sklearn.tree import export_text`
`r =export_text(classifier, feature_names=feature_names)`
`print(r)`

```

| --- alcohol <= 2.50
| | ----olatile acidity <= 0.50
| | | --- sulphates <= 1.50
| | | | ---density <= 0.50
| | | | | --- pH <= 1.50
| | | | | | --- sulphates <= 0.50
| | | | | | | --- class: 2
| | | | | | |--- sulphates > 0.50
| | | | | | |--- class: 7
| | | | | --- pH > 1.50
| | | | | |--- class: 2
| | | | |--- density > 0.50
| | | | |--- class: 2
| | | --- sulphates > 1.50

```

```

class: 2
    --- free sulfur dioxide > 0.50
    |--- class: 7
    --- free sulfur dioxide > 1.50
    |--- pH <= 0.50
    |--- residual sugar <= 2.50
    |--- class: 7
    |--- residual sugar > 2.50
    |--- class: 2
    --- pH > 0.50
    |--- class: 2

```

In [78]:

```
# K FOLD CROSS VALIDATION, K=3

from sklearn.model_selection import KFold
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

X = dt.iloc[:, :-1]
y = dt.iloc[:, -1]

#Implementing cross validation

k=3
kf = KFold(n_splits=k, random_state=None)
model = LogisticRegression(solver= 'liblinear')

acc_score = []

for train_index , test_index in kf.split(X):
    X_train,X_test=X.iloc[train_index,:],X.iloc[test_index,:]
    y_train , y_test = y[train_index] , y[test_index]

    model.fit(X_train,y_train)
    pred_values = model.predict(X_test)

    acc = accuracy_score(pred_values , y_test)
    acc_score.append(acc)

avg_acc_score = sum(acc_score)/k

print('accuracy of each fold - {}'.format(acc_score))
print('Avg accuracy : {}'.format(avg_acc_score))
```

```
accuracy of each fold - [0.874296435272045, 0.8667917448405253, 0.8780487804878049]
Avg accuracy : 0.8730456535334584
```